



**COMP9321**

# **Data Services Engineering**


**Term 1, 2019**

**Week 5 Lecture 2**

# Unsupervised Learning

COMP9321 2019T1

# Supervised learning



Supervised learning: given labeled examples

# Unsupervised learning



Unupervised learning: given data, i.e. examples, but no labels

# Unsupervised Learning

Definition of Unsupervised Learning:

Learning useful structure *without* labeled classes, optimization criterion, feedback signal, or any other information beyond the raw data

# Unsupervised learning applications

learn clusters/groups without any label

customer segmentation (i.e. grouping)

image compression

bioinformatics: learn motifs

find important features

...

# Today's Agenda

- K-Means
- Apriori

A scatter plot on a white background with a yellow footer bar. It contains numerous small grey squares representing data points. There are four distinct clusters of these points. Each cluster has a corresponding colored triangle representing a centroid: a blue triangle in the top-left cluster, a green triangle in the top-right cluster, an orange triangle in the bottom-left cluster, and a yellow triangle in the bottom-right cluster. The text 'Clustering' and 'Unsupervised Learning' is centered in the middle of the plot area.

# Clustering

Unsupervised Learning



# Clustering

Clustering:

- Unsupervised learning
- Requires data, but no labels
- Detect patterns


# Motivations of Clustering

- exploratory data analysis
  - understanding general characteristics of data
  - visualizing data
- generalization – infer something about an instance (e.g. a gene) based on how it relates to other instances

# Paradigms


## Flat algorithms

- Usually start with a random (partial) partitioning
- Refine it iteratively
  - $K$  means clustering
  - Model based clustering
- Spectral clustering



## Hierarchical algorithms

- Bottom-up, agglomerative
- Top-down, divisive



# Paradigms

Hard clustering: Each example belongs to exactly one cluster

Soft clustering: An example can belong to more than one cluster (probabilistic)

- Makes more sense for applications like creating browsable hierarchies
- You may want to put a pair of sneakers in two clusters: (i) sports apparel and (ii) shoes

Disjunctive (an instance can belong to multiple clusters) vs. non-disjunctive

Deterministic (same clusters produced every time for a given data set) vs. stochastic

# Clustering: Image Segmentation


Break up the image into meaningful or perceptually similar regions




# Clustering: Edge Detection



# Basic Idea of Clustering



# Basic Idea of Clustering





# Basic Idea of Clustering

Group together similar data points (instances)

- How to measure the similarity?
- ✓ What could similar mean?
- How many clusters do we need?

# K-means

Most well-known and popular clustering algorithm:


Step 1. Start with some initial cluster centers (k random points)

Step 2. Iterate:


- Assign/cluster each example to closest center
- Recalculate and change centers as the mean of the points in the cluster.

Step 3. Stop when no points' assignments change


# K-means: an example




# K-means: Initialize centers randomly




# K-means: assign points to nearest center




# K-means: readjust centers



# K-means: assign points to nearest center




# K-means: readjust centers







# K-means: assign points to nearest center



# K-means: readjust centers



# K-means: assign points to nearest center




No changes: Done

# K-means

Iterate:

- **Assign/cluster each example to closest center**
- Recalculate centers as the mean of the points in a cluster




How do we do this?

# K-means

Iterate:

- **Assign/cluster each example to closest center**  
iterate over each point:
  - get distance to each cluster center
  - assign to closest center (hard cluster)
- Recalculate centers as the mean of the points in a cluster



# K-means


Iterate:

- **Assign/cluster each example to closest center**

iterate over each point:

- get **distance** to each cluster center
- assign to closest center (hard cluster)

- Recalculate centers as the mean of the points in a cluster



What distance measure should we use?

# Distance measures

Euclidean:


$$d(x, y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

good for spatial data

# K-means

Iterate:

- Assign/cluster each example to closest center
- Recalculate centers as the mean of the points in a cluster




Where are the cluster centers?



# K-means

Iterate:

- Assign/cluster each example to closest center
- Recalculate centers as the mean of the points in a cluster




# K-means

Iterate:

- Assign/cluster each example to closest center
- Recalculate centers as the mean of the points in a cluster

e.g., for a set of instances that have been assigned to a cluster  $c_j$ , we recompute the mean of the cluster as follow

$$\mu(c_j) = \frac{\sum_{\vec{x}_i \in c_j} \vec{x}_i}{|c_j|}$$



# K-means

given : a set  $X = \{\vec{x}_1 \dots \vec{x}_n\}$  of instances

select  $k$  initial cluster centers  $\vec{f}_1 \dots \vec{f}_k$

while stopping criterion not true do

for all clusters  $c_j$  do

// determine which instances are assigned to this cluster

$$c_j = \left\{ \vec{x}_i \mid \forall f_l \text{ dist}(\vec{x}_i, \vec{f}_j) < \text{dist}(\vec{x}_i, \vec{f}_l) \right\}$$


for all means  $\vec{f}_j$  do

// update the cluster center

$$\vec{f}_j = \mu(c_j)$$


# Run an example together ~~

Initialization: 4 points, 2 clusters and distance function




$$\begin{aligned} \text{dist}(x_1, f_1) &= 2, & \text{dist}(x_1, f_2) &= 5 \\ \text{dist}(x_2, f_1) &= 2, & \text{dist}(x_2, f_2) &= 3 \\ \text{dist}(x_3, f_1) &= 3, & \text{dist}(x_3, f_2) &= 2 \\ \text{dist}(x_4, f_1) &= 11, & \text{dist}(x_4, f_2) &= 6 \end{aligned}$$


$$\text{dist}(x_i, x_j) = \sum_e |x_{i,e} - x_{j,e}|$$



$$\begin{aligned} f_1 &= \left\langle \frac{4+4}{2}, \frac{1+3}{2} \right\rangle = \langle 4, 2 \rangle \\ f_2 &= \left\langle \frac{6+8}{2}, \frac{2+8}{2} \right\rangle = \langle 7, 5 \rangle \end{aligned}$$



$$\begin{aligned} \text{dist}(x_1, f_1) &= 1, & \text{dist}(x_1, f_2) &= 7 \\ \text{dist}(x_2, f_1) &= 1, & \text{dist}(x_2, f_2) &= 5 \\ \text{dist}(x_3, f_1) &= 2, & \text{dist}(x_3, f_2) &= 4 \\ \text{dist}(x_4, f_1) &= 10, & \text{dist}(x_4, f_2) &= 4 \end{aligned}$$



$$\begin{aligned} f_1 &= \left\langle \frac{4+4+6}{3}, \frac{1+3+2}{3} \right\rangle = \langle 4.67, 2 \rangle \\ f_2 &= \left\langle \frac{8}{1}, \frac{8}{1} \right\rangle = \langle 8, 8 \rangle \end{aligned}$$

# K-means loss function

K-means tries to minimize what is called the “k-means” loss function:

$$loss = \sum_{i=1}^n d(x_i, \mu_k)^2 \quad \text{where } \mu_k \text{ is cluster center for } x_i$$

that is, the sum of the squared distances from each point to the associated cluster center

# Minimizing k-means loss

Iterate:

1. Assign/cluster each example to closest center
  2. Recalculate centers as the mean of the points in a cluster
- 

$$loss = \sum_{i=1}^n d(x_i, \mu_k)^2 \quad \text{where } \mu_k \text{ is cluster center for } x_i$$

Does each step of k-means move towards reducing this loss function (or at least not increasing)?

# Minimizing k-means loss

Iterate:

1. Assign/cluster each example to closest center
  2. Recalculate centers as the mean of the points in a cluster
- 

$$loss = \sum_{i=1}^n d(x_i, \mu_k)^2 \quad \text{where } \mu_k \text{ is cluster center for } x_i$$

This isn't quite a complete proof/argument, but:

1. Any other assignment would end up in a larger loss
2. The mean of a set of values minimizes the squared error

# Minimizing k-means loss

Iterate:

1. Assign/cluster each example to closest center
  2. Recalculate centers as the mean of the points in a cluster
- 

$$loss = \sum_{i=1}^n d(x_i, \mu_k)^2 \quad \text{where } \mu_k \text{ is cluster center for } x_i$$

Does this mean that k-means will always find the minimum loss/clustering?



# Minimizing k-means loss

Iterate:

1. Assign/cluster each example to closest center
  2. Recalculate centers as the mean of the points in a cluster
- 

$$loss = \sum_{i=1}^n d(x_i, \mu_k)^2 \quad \text{where } \mu_k \text{ is cluster center for } x_i$$

NO! It will find *a minimum*.

Unfortunately, the k-means loss function is generally not convex and for most problems has many, many minima

We're only guaranteed to find one of them

# Properties of K-means

Guaranteed to converge in a finite number of iterations

Running time per iteration

1. Assign data points to closest cluster center  $O(KN)$  time
2. Change the cluster center to the average of its assigned points  $O(N)$

# K-means variations/parameters

Start with some initial cluster centers

Iterate:

- Assign/cluster each example to closest center
- Recalculate centers as the mean of the points in a cluster

What are some other variations/parameters we haven't specified?

# K-means variations/parameters


Initial (seed) cluster centers

## Convergence

- A fixed number of iterations
- partitions unchanged
- Cluster centers don't change

K!

# K-means: Initialize centers randomly



What would happen here?

Seed selection ideas?

# Seed choice

Results can vary drastically based on random seed selection

Some seeds can result in poor convergence rate, or convergence to sub-optimal clusterings

## Common heuristics

- Random centers in the space
- Randomly pick examples
- Points least similar to any existing center (furthest centers heuristic)
- **Try out multiple starting points**
- Initialize with the results of another clustering method

# Furthest centers heuristic

$\mu_1$  = pick random point

for  $i = 2$  to  $K$ :

$\mu_i$  = point that is furthest from **any** previous centers


---

$$\mu_i = \underbrace{\arg \max_x}_{\text{point with the largest distance to any previous center}} \underbrace{\min_{\mu_j : 1 < j < i} d(x, \mu_j)}_{\text{smallest distance from } x \text{ to any previous center}}$$

point with the largest  
distance to any previous  
center

smallest distance from  $x$  to  
any previous center


# K-means: Initialize furthest from centers



Pick a random point for the first center




# K-means: Initialize furthest from centers



What point will be chosen next?


# K-means: Initialize furthest from centers



Furthest point from center

What point will be chosen next?


# K-means: Initialize furthest from centers



Furthest point from center

What point will be chosen next?


# K-means: Initialize furthest from centers



Furthest point from center

Any issues/concerns with this approach?


# Furthest points concerns



If  $k = 4$ , which points will get chosen?




# Furthest points concerns



If we do a number of trials, will we get different centers?

# Furthest points concerns



Doesn't deal well with outliers

# K-means++

$\mu_1$  = pick random point

for  $k = 2$  to **K**:

for  $i = 1$  to **N**:

$s_i = \min d(x_i, \mu_{1\dots k-1})$  // smallest distance to any center

$\mu_k$  = randomly pick point *proportionate* to *s*

How does this help?



# K-means++

$\mu_1$  = pick random point

for  $k = 2$  to  $K$ :

for  $i = 1$  to  $N$ :

$s_i = \min d(x_i, \mu_{1\dots k-1})$  // smallest distance to any center

$\mu_k$  = randomly pick point *proportionate* to *s*

- Makes it possible to select other points
  - if #points  $\gg$  #outliers, we will pick good points
- Makes it non-deterministic, which will help with random runs
- Nice theoretical guarantees!

# What Is A Good Clustering?

Internal criterion: A good clustering will produce high quality clusters in which:

- the intra-class (that is, intra-cluster) similarity is high
- the inter-class similarity is low
- The measured quality of a clustering depends on both the document representation and the similarity measure used

# Clustering Evaluation

- Intra-cluster cohesion (compactness):
  - Cohesion measures how near the data points in a cluster are to the cluster centroid.
  - Sum of squared error (SSE) is a commonly used measure.
- Inter-cluster separation (isolation):
  - Separation means that different cluster centroids should be far away from one another.
- In most applications, expert judgments are still the key

# Association rule

Unsupervised Learning

# Mining Association Rules

- What is Association rule mining
- Apriori Algorithm
- Additional Measures of rule interestingness

# What Is Association Rule Mining?

- Finding frequent patterns, associations, correlations, or causal structures among sets of items in transaction databases
- Understand customer buying habits by finding associations and correlations between the different items that customers place in their “shopping basket”
- Applications
  - Basket data analysis, cross-marketing, catalog design, loss-leader analysis, web log analysis, fraud detection (supervisor->examiner)

# What Is Association Rule Mining?

Rule form

Antecedent  $\rightarrow$  Consequent [support, confidence]

*(support and confidence are user defined measures of interestingness)*

## Examples

buys(x, “computer”)  $\rightarrow$  buys(x, “financial management software”) [0.5%, 60%]

age(x, “30..39”)  $\wedge$  income(x, “42..48K”)  $\rightarrow$  buys(x, “car”) [1%, 75%]

# How can Association Rules be used?

## Stories – Beer and Diapers



- ♦ Diapers and Beer. Most famous example of market basket analysis for the last few years. If you buy diapers, you tend to buy beer.
- T. Blischok headed Terradata's Industry Consulting group.
- K. Heath ran self joins in SQL (1990), trying to find two itemsets that have baby items, which are particularly profitable.
- Found this pattern in their data of 50 stores/90 day period.
- Unlikely to be significant, but it's a nice example that explains associations well.

*Ronny Kohavi ICML 1998*

Probably mom was calling dad at work to buy diapers on way home and he decided to buy a six-pack as well.

The retailer could move diapers and beers to separate places and position high profit items of interest to young fathers along the path.



# How can Association Rules be used?

Let the rule discovered be

$$\{\text{Bagels}, \dots\} \rightarrow \{\text{Potato Chips}\}$$

**Potato chips as consequent** → Can be used to determine what should be done to boost its sales

**Bagels in the antecedent** → Can be used to see which products would be affected if the store discontinues selling bagels

**Bagels in antecedent and Potato chips in the consequent** → Can be used to see what products should be sold with Bagels to promote sale of Potato Chips

# Association Rule: Basic Concepts

Given:

- (1) database of *transactions*,
- (2) each transaction is a *list of items* purchased by a customer in a visit

Find:

- all rules that correlate the presence of one set of items (itemset) with that of another set of items
- E.g., 98% of people who purchase tires and auto accessories also get automotive services done

# Rule basic Measures:

## Support and Confidence

$$A \rightarrow B [s, c]$$

**Support:** denotes the frequency of the rule within transactions. A high value means that the rule involve a great part of database.

$$\text{support}(A \rightarrow B [s, c]) = p(A \rightarrow B)$$

**Confidence:** denotes the percentage of transactions containing A which contain also B. It is an estimation of conditioned probability .

$$\text{confidence}(A \rightarrow B [s, c]) = p(B|A) = \text{sup}(A,B)/\text{sup}(A).$$

# Example

Trans. Id	Purchased Items
1	A,D
2	A,C
3	A,B,C
4	B,E,F

Itemset:

A,B or B,E,F

Support of an itemset:

$\text{Sup}(A,B)=1$     $\text{Sup}(A,C)=2$

Frequent pattern:

Given min. sup=2,

{A,C} is a frequent pattern

For minimum support = 50% and minimum confidence = 50%,  
we have the following rules

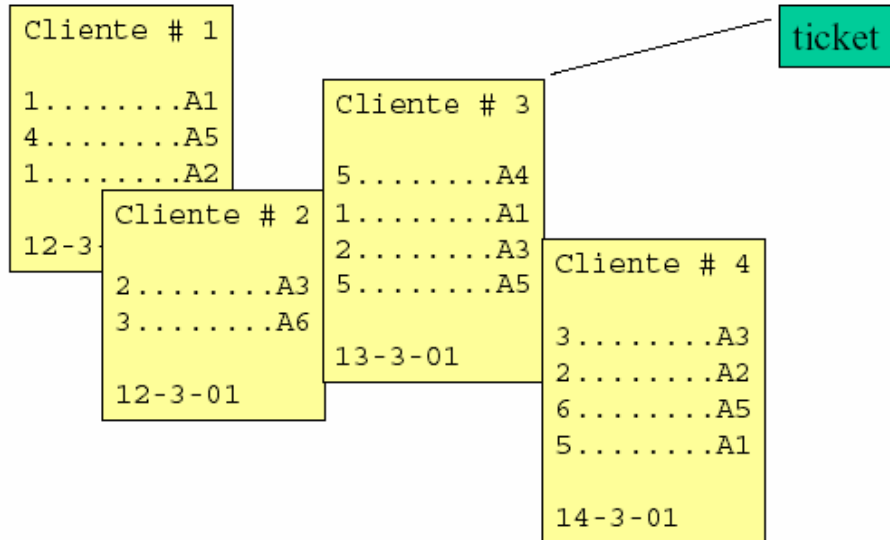
$A \rightarrow C$       with 50% support and      66% confidence

$C \rightarrow A$       with 50% support and      100% confidence

# Mining Association Rules

- What is Association rule mining
- Apriori Algorithm

# Boolean association rules



Each transaction is represented by a Boolean vector

Cliente	A1	A2	A3	A4	A5	A6	A7	A8	A9	A10	A11	A12	A13
1	1	1	0	0	1	0	0	0	0	0	0	0	0
2	0	0	1	0	0	1	0	0	0	0	0	0	0
3	1	0	1	1	1	0	0	0	0	0	0	0	0
4	1	1	1	0	1	0	0	0	0	0	0	0	0
5	0	0	1	0	0	1	0	1	1	1	0	0	0
6	0	1	0	0	0	0	0	1	0	1	0	0	0
7	1	0	0	0	0	0	1	1	0	1	0	1	1
8	0	1	0	0	0	0	0	1	0	0	0	0	0
9	0	0	0	0	0	0	0	0	0	1	0	1	0

# Mining Association Rules - An Example

Transaction ID	Items Bought
2000	A,B,C
1000	A,C
4000	A,D
5000	B,E,F

Min. support 50%  
Min. confidence 50%

Frequent Itemset	Support
{A}	75%
{B}	50%
{C}	50%
{A,C}	50%

For rule  $A \rightarrow C$  :

support = support({A , C }) = 50%

confidence = support({A , C }) / support({A }) = 66.6%

# The Apriori principle

**Any subset of a frequent itemset  
must be frequent**

A transaction containing {beer, diaper, nuts} also contains {beer, diaper}

{beer, diaper, nuts} is frequent

→ {beer, diaper} must also be frequent

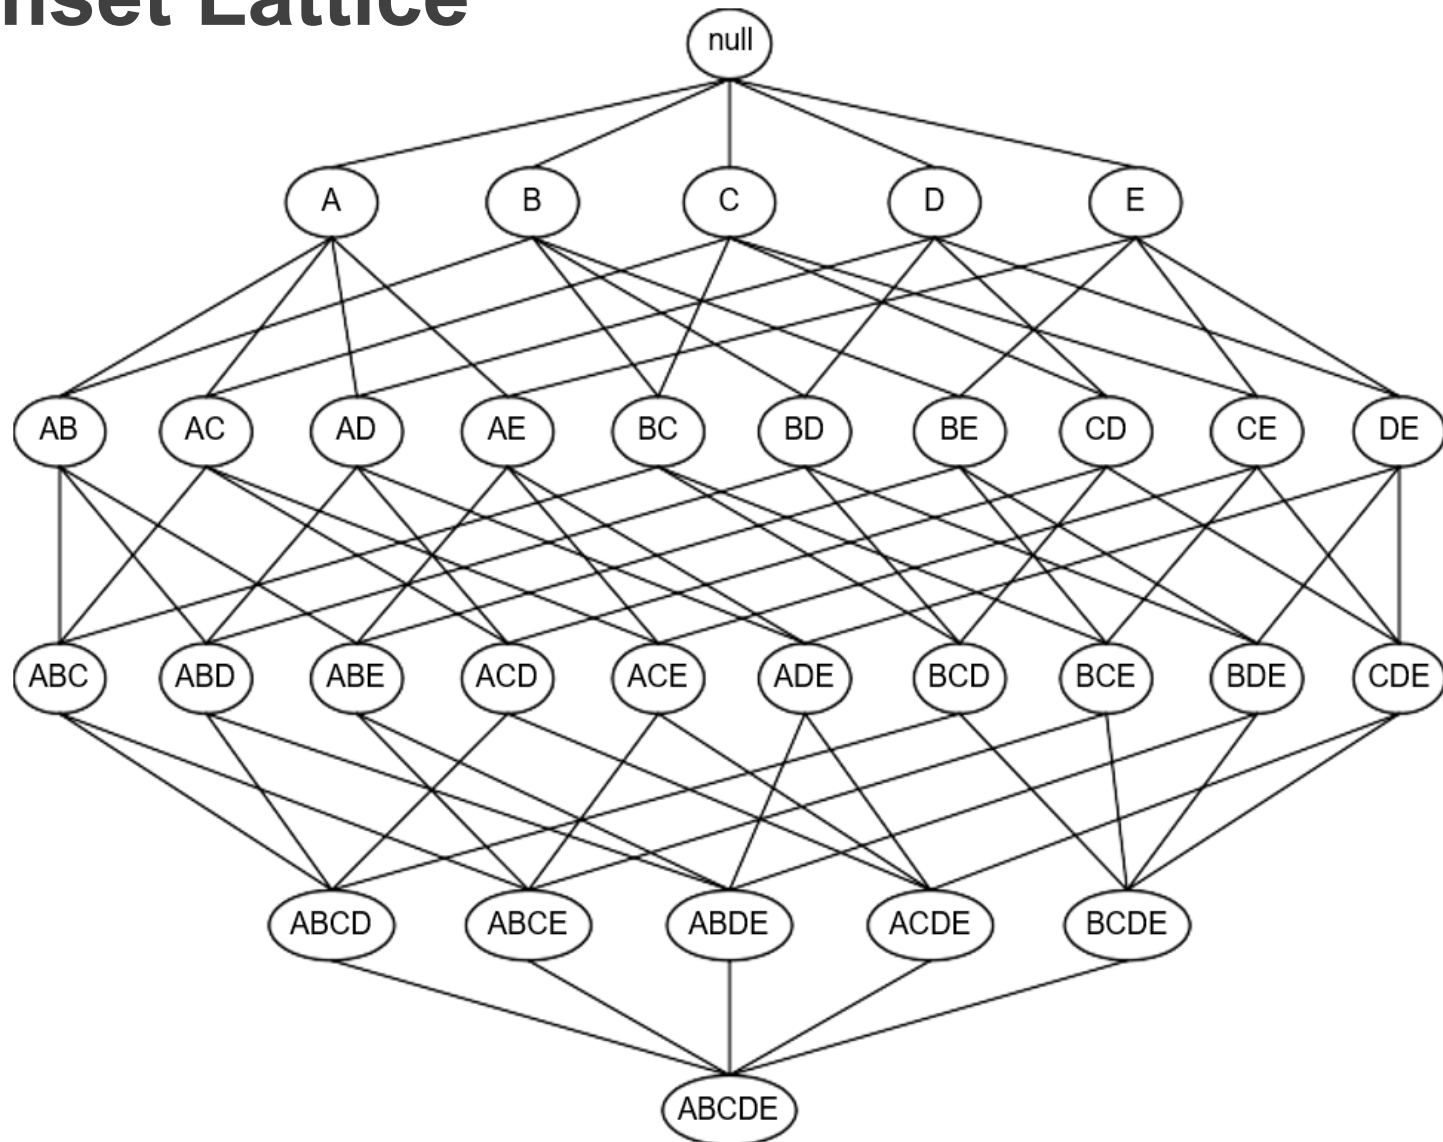


# The Apriori principle

No superset of any infrequent itemset  
should be generated or tested

Many item combinations can be pruned

# Itemset Lattice

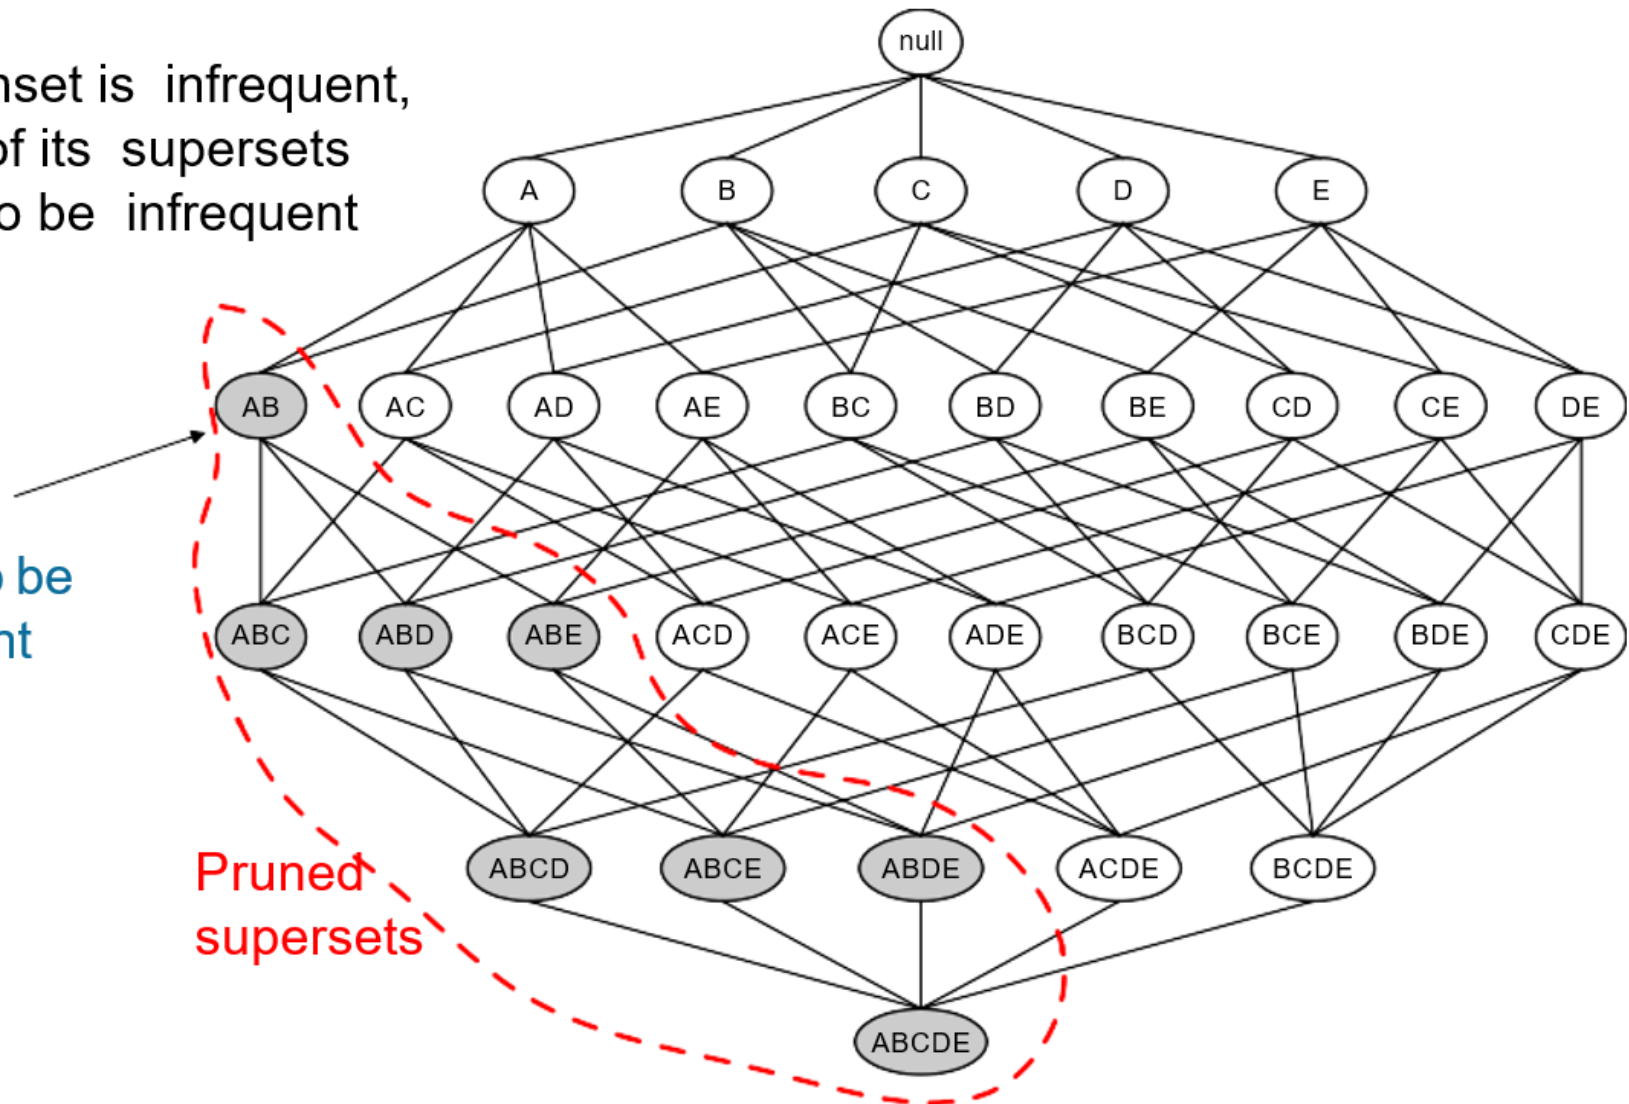


# Apriori principle for pruning candidates

If an itemset is infrequent,  
then all of its supersets  
must also be infrequent

Found to be  
Infrequent

Pruned  
supersets



# Mining Frequent Itemsets (the Key Step)

**Find the *frequent itemsets*:** the sets of items that have minimum support

A subset of a frequent itemset must also be a frequent itemset

Generate length  $(k+1)$  candidate itemsets from length  $k$  frequent itemsets, and

Test the candidates against DB to determine which are in fact frequent

Use the **frequent itemsets to generate association rules.**

Generation is straightforward

# The Apriori Algorithm — Example

Database D

TID	Items
100	1 3 4
200	2 3 5
300	1 2 3 5
400	2 5

Scan D

$C_1$

itemset	sup.
{1}	2
{2}	3
{3}	3
{4}	1
{5}	3

Min. support: 2 transactions

$L_1$

itemset	sup.
{1}	2
{2}	3
{3}	3
{5}	3

$L_2$

itemset	sup
{1 3}	2
{2 3}	2
{2 5}	3
{3 5}	2

$C_2$

itemset	sup
{1 2}	1
{1 3}	2
{1 5}	1
{2 3}	2
{2 5}	3
{3 5}	2

Scan D

$C_2$

item set
{1 2}
{1 3}
{1 5}
{2 3}
{2 5}
{3 5}

$C_3$

itemset
{2 3 5}

Scan D

$L_3$

itemset	sup
{2 3 5}	2

# How to Generate Candidates?

The items in  $L_{k-1}$  are listed in an order

Step 1: self-joining  $L_{k-1}$

- insert into  $C_k$
- select  $p.item_1, p.item_2, \dots, p.item_{k-1}, q.item_{k-1}$
- from  $L_{k-1} p, L_{k-1} q$
- where  $p.item_1 = q.item_1, \dots, p.item_{k-2} = q.item_{k-2},$   
 $p.item_{k-1} < q.item_{k-1}$


A	D	E
 A	 D	^ F

## Step 2: pruning

for all *itemsets*  $c$  in  $C_k$  do

for all  $(k-1)$ -subsets  $s$  of  $c$  do

if ( $s$  is not in  $L_{k-1}$ ) then delete  $c$  from  $C_k$



# Example of Generating Candidates

$L_3 = \{abc, abd, acd, ace, bcd\}$

**Self-joining**:  $L_3 * L_3$

$abcd$  from  $abc$  and  $abd$

$acde$  from  $acd$  and  $ace$

**Pruning** (*before counting its support*):

$acde$  is removed because  $ade$  is not in  $L_3$

$C_4 = \{abcd\}$



# The Apriori Algorithm

$C_k$ : Candidate itemset of size  $k$      $L_k$ : frequent itemset of size  $k$

**Join Step:**  $C_k$  is generated by joining  $L_{k-1}$  with itself

**Prune Step:** Any  $(k-1)$ -itemset that is not frequent cannot be a subset of a frequent  $k$ -itemset

**Algorithm:**

$L_1 = \{\text{frequent items}\};$

**for** ( $k = 1; L_k \neq \emptyset; k++$ ) **do begin**

$C_{k+1}$  = candidates generated from  $L_k$ ;

**for each** transaction  $t$  in database **do**

        increment the count of all candidates in  $C_{k+1}$  that are contained in  $t$

$L_{k+1}$  = candidates in  $C_{k+1}$  with min\_support

**end**

**return**  $L = \cup_k L_k$ ;

# Generating AR from frequent itemsets

$$\text{Confidence } (A \rightarrow B) = P(B|A) = \frac{\text{support\_count}(\{A,B\})}{\text{support\_count}(\{A\})}$$

For every frequent itemset  $x$ , generate all non-empty subsets of  $x$

For every non-empty subset  $s$  of  $x$ , output the rule  
“ $s \rightarrow (x-s)$ ” if

$$\frac{\text{support\_count}(\{x\})}{\text{support\_count}(\{s\})} \geq \text{min\_conf}$$

# Generating AR from frequent itemsets

Given a frequent itemset  $L$ , find all non-empty subsets  $f \subset L$  such that  $f \rightarrow L - f$  satisfies the minimum confidence requirement

If  $\{A,B,C,D\}$  is a frequent itemset, candidate rules:

$ABC \rightarrow D, ABD \rightarrow C, ACD \rightarrow B, BCD \rightarrow A, A \rightarrow BCD, B \rightarrow ACD,$   
 $C \rightarrow ABD, D \rightarrow ABC, AB \rightarrow CD, AC \rightarrow BD, AD \rightarrow BC, BC \rightarrow AD,$   
 $BD \rightarrow AC, CD \rightarrow AB$

If  $|L| = k$ , then there are  $2^k - 2$  candidate association rules (ignoring  $L \rightarrow \emptyset$  and  $\emptyset \rightarrow L$ )

# From Frequent Itemsets to Association Rules

*Q: Given frequent set  $\{A,B,E\}$ , what are possible association rules?*

$A \rightarrow B, E$

$A, B \rightarrow E$

$A, E \rightarrow B$

$B \rightarrow A, E$

$B, E \rightarrow A$


$E \rightarrow A, B$

$\_\_\_ \rightarrow A, B, E$  (empty rule), or  $\text{true} \rightarrow A, B, E$


# Generating Rules: example

Trans-ID	Items
1	ACD
2	BCE
3	ABCE
4	BE
5	ABCE

Min\_support: 60%  
Min\_confidence: 75%



Frequent Itemset	Support
{ <b>BCE</b> },{AC}	60%
{BC},{CE},{A}	60%
{BE},{B},{C},{E}	80%



Rule	Conf.
{BC} => {E}	100%
{BE} => {C}	75%
{CE} => {B}	100%
{B} => {CE}	75%
{C} => {BE}	75%
{E} => {BC}	75%

# Database TDB

Tid	Items
10	A, C, D
20	B, C, E
30	A, B, C, E
40	B, E

1<sup>st</sup> scan

$C_1$

Itemset	sup
{A}	2
{B}	3
{C}	3
{D}	1
{E}	3

$L_1$

Itemset	sup
{A}	2
{B}	3
{C}	3
{E}	3

$L_2$

Itemset	sup
{A, C}	2
{B, C}	2
{B, E}	3
{C, E}	2

$C_2$

Itemset	sup
{A, B}	1
{A, C}	2
{A, E}	1
{B, C}	2
{B, E}	3
{C, E}	2

2<sup>nd</sup> scan

$C_2$

Itemset	sup
{A, B}	1
{A, C}	2
{A, E}	1
{B, C}	2
{B, E}	3
{C, E}	2

$C_3$

Itemset	sup
{B, C, E}	2

3<sup>rd</sup> scan

$L_3$

Itemset	sup
{B, C, E}	2

# Challenges of Frequent Pattern Mining

## Challenges

- Multiple scans of transaction database

- Huge number of candidates

- Tedious workload of support counting for candidates

## Improving Apriori: general ideas

- Reduce number of transaction database scans

- Shrink number of candidates

- Facilitate support counting of candidates